

Introduction to Theano

Arnaud Bergeron

(slides adapted by Frédéric Bastien from slides by Ian G.)

(further adapted by Arnaud Bergeron)

February 26, 2015

High level

- ▶ Overview of library (3 min)
- ▶ Building expressions (30 min)
- ▶ Compiling and running expressions (30 min)
- ▶ Modifying expressions (25 min)
- ▶ Debugging (30 min)
- ▶ Citing Theano (2 min)

Overview of Library

Theano is many things

- ▶ Language
- ▶ Compiler
- ▶ Python library

Theano language:

- ▶ Operations on scalar, vector, matrix, tensor, and sparse variables
- ▶ Linear algebra
- ▶ Element-wise nonlinearities
- ▶ Convolution
- ▶ Extensible

Overview

Using Theano:

- ▶ define expression $f(x, y) = x + y$

```
>>> z = x + y
```

- ▶ compile expression

```
>>> f = theano.function([x, y], z)
```

- ▶ execute expression

```
>>> f(1, 2)
3
```

Building expressions

- ▶ Scalars
- ▶ Vectors
- ▶ Matrices
- ▶ Tensors
- ▶ Reduction
- ▶ Dimshuffle

Scalar math

```
from theano import tensor as T
x = T.scalar()
y = T.scalar()
z = x+y
w = z*x
a = T.sqrt(w)
b = T.exp(a)
c = a ** b
d = T.log(c)
```

Vector math

```
from theano import tensor as T
x = T.vector()
y = T.vector()
# Scalar math applied elementwise
a = x * y
# Vector dot product
b = T.dot(x, y)
# Broadcasting
c = a + b
```


Matrix math

```
from theano import tensor as T
x = T.matrix()
y = T.matrix()
a = T.vector()
# Matrix-matrix product
b = T.dot(x, y)
# Matrix-vector product
c = T.dot(x, a)
```

Tensors

- ▶ Dimensionality defined by length of “broadcastable” argument
- ▶ Can add (or do other elemwise op) on two tensors with same dimensionality
- ▶ Duplicate tensors along broadcastable axes to make size match

```
from theano import tensor as T
tensor3 = T.TensorType(
    broadcastable=(False, False, False),
    dtype='float32')
x = tensor3()
```

Reductions

```
from theano import tensor as T
tensor3 = T.TensorType(
    broadcastable=(False, False, False),
    dtype='float32')
x = tensor3()
total = x.sum()
marginals = x.sum(axis=(0, 2))
mx = x.max(axis=1)
```

Dimshuffle

```
from theano import tensor as T
tensor3 = T.TensorType(
    broadcastable=(False, False, False),
    dtype='float32')
x = tensor3()
y = x.dimshuffle((2, 1, 0))
a = T.matrix()
b = a.T
# Same as b
c = a.dimshuffle((0, 1))
# Adding to larger tensor
d = a.dimshuffle((0, 1, 'x'))
e = a + d
```

Work through the "Building Expressions" section of the ipython notebook.

Compiling and running expression

- ▶ `theano.function`
- ▶ shared variables and updates
- ▶ compilation modes
- ▶ compilation for GPU
- ▶ optimizations

```
>>> from theano import tensor as T
>>> x = T.scalar()
>>> y = T.scalar()
>>> from theano import function
>>> # first arg is list of symbolic inputs
>>> # second arg is symbolic output
>>> f = function([x, y], x + y)
>>> # Call it with numerical values
>>> # Get a numerical output
>>> f(1., 2.)
array(3.0)
```

Shared variables

- ▶ It's hard to do much with purely functional programming
- ▶ *shared variables* add just a little bit of imperative programming
- ▶ A *shared variable* is a buffer that stores a numerical value for a Theano variable
- ▶ Can write to as many shared variables as you want, once each, at the end of the function
- ▶ Modify outside Theano function with `get_value()` and `set_value()` methods.

Shared variable example

```
>>> from theano import shared
>>> x = shared(0.)
# Can also use a dict for more complex code
>>> updates = [(x, x + 1)]
>>> f = function([], updates=updates)
>>> f()
>>> x.get_value()
1.0
>>> x.set_value(100.)
>>> f()
>>> x.get_value()
101.0
```

Which dict?

- ▶ Use `theano.compat.python2x.OrderedDict`
- ▶ Not `collections.OrderedDict`
 - ▶ This isn't available in older versions of python, and will limit the portability of your code.
- ▶ Not `{}` aka dict
 - ▶ The iteration order of this built-in class is not deterministic so if Theano accepted this, the same script could compile different C programs each time you run it.

Compilation modes

- ▶ Can compile in different modes to get different kinds of programs
- ▶ Can specify these modes very precisely with arguments to `theano.function()`
- ▶ Can use a few quick presets with environment variable flags

Example preset compilation modes

FAST_RUN Default. Spends a lot of time on compilation to get an executable that runs fast.

FAST_COMPILE Doesn't spend much time compiling. Executable usually uses python instead of compiled C code. Runs slow.

DEBUG_MODE Adds lots of checks. Raises error messages in situations other modes don't check for.

Compilation for GPU

- ▶ Theano's current back-end only supports 32 bit on GPU
- ▶ CUDA supports 64 bit, but is slow in gamer card
- ▶ `T.fscalar`, `T.fvector`, `T.fmatrix` are all 32 bit
- ▶ `T.scalar`, `T.vector`, `T.matrix` resolve to 32 or 64 bit depending on theano's floatX flag
- ▶ floatX is float64 by default, set it to float32
- ▶ Set the device flag to gpu (or a specific gpu, like gpu0)
- ▶ Optional: `warn_float64={'ignore', 'warn', 'raise', 'pdb'}`

Optimizations

- ▶ Theano changes the symbolic expressions you write before converting them to C code
- ▶ It makes them faster
 - ▶ $(x + y) + (x + y) \rightarrow 2 \times (x + y)$
- ▶ It makes them more stable
 - ▶ $\exp(a) / \sum \exp(a) \rightarrow \text{softmax}(a)$

Optimizations (2)

Sometimes optimizations discard error checking and produce incorrect output rather than an exception.

```
>>> x = T.scalar()
>>> f = function([x], x/x)
>>> f(0.)
array(1.0)
```

Work through the "Compiling and Running" section of the ipython notebook.

Modifying expressions

- ▶ The `grad()` method
- ▶ Variable nodes
- ▶ Types
- ▶ Ops
- ▶ Apply nodes

The `grad()` method

```
>>> x = T.scalar('x')
>>> y = 2. * x
>>> g = T.grad(y, x)
>>> from theano.printing import min_informative_str
# Print the unoptimized graph
>>> print min_informative_str(g)
A. Elemwise{mul}
  B. Elemwise{second,no_inplace}
    C. Elemwise{mul,no_inplace}
      D. TensorConstant{2.0}
        E. x
          F. TensorConstant{1.0}
<D>
```

The `grad()` method

```
>>> x = T.scalar('x')
>>> y = 2. * x
>>> g = T.grad(y, x)
>>> from theano.printing import min_informative_str
# Print the optimized graph
>>> f = theano.function([x], g)
>>> theano.printing.debugprint(f)
DeepCopyOp [@A] '' 0
|TensorConstant{2.0} [@B]
```

Theano variables

- ▶ A *variable* is a theano expression.
- ▶ Can come from `T.scalar()`, `T.matrix()`, etc.
- ▶ Can come from doing operations on other variables.
- ▶ Every variable has a `type` field, identifying its *type*, such as `TensorType((True, False), 'float32')`
- ▶ Variables can be thought of as nodes in a graph

- ▶ An Op is any class that describes a function operating on some variables
- ▶ Can call the op on some variables to get a new variable or variables
- ▶ An Op class can supply other forms of information about the function, such as its derivative

Apply nodes

- ▶ The Apply class is a specific instance of an application of an Op.
- ▶ Notable fields:
 - `op` The Op to be applied
 - `inputs` The Variables to be used as input
 - `outputs` The Variables produced
- ▶ The `owner` field on variables identifies the Apply that created it.
- ▶ Variable and Apply instances are nodes and owner/inputs/outputs identify edges in a Theano graph.

Work through the "Modifying" section in the ipython notebook.

- ▶ `DEBUG_MODE`
- ▶ Error message
- ▶ `theano.printing.debugprint()`
- ▶ `min_informative_str()`
- ▶ `compute_test_value`
- ▶ Accessing the `FunctionGraph`

Error message: code

```
import numpy as np
import theano
import theano.tensor as T
x = T.vector()
y = T.vector()
z = x + x
z = z + y
f = theano.function([x, y], z)
f(np.ones((2,)), np.ones((3,)))
```

Error message I

```
Traceback (most recent call last):
  File "test.py", line 9, in <module>
    f(np.ones((2,)), np.ones((3,)))
  File "/Users/anakha/Library/Python/2.7/site-packages
    /theano/compile/function_module.py", line 606,
    in __call__
    storage_map=self.fn.storage_map)
  File "/Users/anakha/Library/Python/2.7/site-packages
    /theano/compile/function_module.py", line 595,
    in __call__
    outputs = self.fn()
ValueError: Input dimension mis-match. (input[0].shape
  [0] = 3, input[1].shape[0] = 2)
Apply node that caused the error: Elemwise{add,
  no_inplace}(<TensorType(float64, vector)>, <
  TensorType(float64, vector)>, <TensorType(float64,
  vector)>)
```

Error message II

```
Inputs types: [TensorType(float64, vector), TensorType
               (float64, vector), TensorType(float64, vector)]
Inputs shapes: [(3,), (2,), (2,)]
Inputs strides: [(8,), (8,), (8,)]
Inputs values: [array([ 1.,  1.,  1.]), array([ 1.,
        1.]), array([ 1.,  1.])]
```

HINT: Re-running with most Theano optimization disabled could give you a back-trace of when this node was created. This can be done with by setting the Theano flag 'optimizer=fast_compile'. If that does not work, Theano optimizations can be disabled with 'optimizer=None'.

HINT: Use the Theano flag 'exception_verbosity=high' for a debugprint and storage map footprint of this apply node.

Error message: exception_verbosity=high

Debugprint of the apply node:

```
Elemwise{add,no_inplace} [@A] <TensorType(float64,
    vector)> ''
|<TensorType(float64, vector)> [@B] <TensorType(
    float64, vector)>
|<TensorType(float64, vector)> [@C] <TensorType(
    float64, vector)>
|<TensorType(float64, vector)> [@C] <TensorType(
    float64, vector)>
```

Storage map footprint:

- <TensorType(float64, vector)>, Shape: (3,),
ElemSize: 8 Byte(s), TotalSize: 24 Byte(s)
- <TensorType(float64, vector)>, Shape: (2,),
ElemSize: 8 Byte(s), TotalSize: 16 Byte(s)

Error message: optimizer=fast_compile

```
Backtrace when the node is created:  
File "test.py", line 7, in <module>  
    z = z + y
```

debugprint

```
>>> from theano.printing import debugprint
>>> debugprint(a)
Elemwise{mul,no_inplace} [@A] ''
|TensorConstant{2.0} [@B]
|Elemwise{add,no_inplace} [@C] 'z'
|<TensorType(float64, scalar)> [@D]
|<TensorType(float64, scalar)> [@E]
```

min_informative_str

```
>>> x = T.scalar()
>>> y = T.scalar()
>>> z = x + y
>>> z.name = 'z'
>>> a = 2. * z
>>> from theano.printing import min_informative_str
>>> print min_informative_str(a)
A. Elemwise{mul,no_inplace}
B. TensorConstant{2.0}
C. z
```

compute_test_value

```
>>> from theano import config
>>> config.compute_test_value = 'raise'
>>> x = T.vector()
>>> import numpy as np
>>> x.tag.test_value = np.ones((2,))
>>> y = T.vector()
>>> y.tag.test_value = np.ones((3,))
>>> x + y
...
ValueError: Input dimension mis-match.
(input[0].shape[0] = 2, input[1].shape[0] = 3)
```


Accessing a function's fgraph

```
>>> x = T.scalar()
>>> y = x / x
>>> f = function([x], y)
>>> debugprint(f maker.fgraph.outputs[0])
DeepCopyOp [@A] ''
|TensorConstant{1.0} [@B]
```

Work through the "Debugging" section of the ipython notebook.

Citing Theano

Please cite both of the following papers in all work that uses Theano:

- ▶ Bastien, Frédéric, Lamblin, Pascal, Pascanu, Razvan, Bergstra, James, Goodfellow, Ian, Bergeron, Arnaud, Bouchard, Nicolas, and Bengio, Yoshua. Theano: new features and speed improvements. Deep Learning and Unsupervised Feature Learning NIPS 2012 Workshop, 2012.
- ▶ Bergstra, James, Breuleux, Olivier, Bastien, Frédéric, Lamblin, Pascal, Pascanu, Razvan, Desjardins, Guillaume, Turian, Joseph, Warde-Farley, David, and Bengio, Yoshua. Theano: a CPU and GPU math expression compiler. In Proceedings of the Python for Scientific Computing Conference (SciPy), June 2010. Oral Presentation.

Example acknowledgments

We would like to thank the developers of Theano
`\citep{bergstra+al:2010-scipy,Bastien-Theano-2012}`. We would
also like to thank NSERC, Compute Canada, and Calcul Québec
for providing computational resources.

Questions?